

# HW 2 Stats Learn, Date: 02/20/25

Steafan Steinocher ID 1554129

The following Question 1 is in regards to this table:

Obs	X1	X2	X3	Y
1	1	2	-1	Blue
2	1	0	3	Blue
3	0	3	0	Red
4	0	2	-1	Red
5	2	0	-1	Blue
6	1	4	1	Red

## Question 1a:

The nearest neighbor to any  $K = 1$  observation is the observation itself so in this case Blue with a misclassification rate of 0%, and this would be the case for each of the other observations.

## Question 1b:

Where  $K = 3$ , there is a misclassification of 1/6.

Obs	X1	X2	X3	Y
1	Blue	Red	Red	Red
2	Blue	Blue	Red	Blue
3	Red	Red	Blue	Red
4	Red	Blue	Red	Red
5	Blue	Blue	Red	Blue
6	Red	Red	Blue	Red

**Question 1c:****Question 1 c i:**

Because our point of interest is (0,0,0) the distance of each observation from our point of interest is simply the magnitude of the observation.

Obs	Distance from (0,0,0)
1	$\sqrt{6}$
2	$\sqrt{10}$
3	3
4	$\sqrt{5}$
5	$\sqrt{5}$
6	$\sqrt{18}$

**Question 1 c ii:**

- For  $K = 1$  the nearest neighbor prediction could be Red or Blue as they are equidistant from the initial point.

**Question 1 c iii:**

- For  $K = 3$  the nearest neighbor prediction would be Blue as the majority of points nearest to our point of interest are Blue.

**Question 1 c iv:**

- For  $K = 5$  the nearest neighbor prediction would be Blue because the majority of points nearest to our point of interest are Blue.

**Question 2:****Question 2 a:**

The validation set approach splits the data set into two parts.

- Training set which is often seen as 70% of the data set.
- Validation set, which is used to evaluate the models performance which is often roughly 30% of the data set.

A disadvantage to the validation set approach is Data inefficiency, not as much data is used to train a model that adopts this approach.

**Question 2 b:**

Leave one out cross validation works as follows:

- Treat that observation as the validation set.
- Train the model on the remaining  $(n-1)$  observations, hence leave one out.
- Evaluate the model's performance on the held-out observation.

One of the advantages of LOOCV is that because more data is used, there is lower variance.

**Question 2 c:**

The main disadvantage of LOOCV that I have seen in this HW is the amount of time for computer. I think on EXTREMELY LARGE data sets, compute time will be long.

**Question 2 d:**

Both validation set approach and LOOCV can be applied to any supervised learning algorithm, not just k-Nearest Neighbors. Like, linear regression.

**Question 3:**

```
#|eval: false
#Needed libraries
library(mlbench)
library(ISLR)
library(caret)#For KNN
```

Loading required package: ggplot2

Loading required package: lattice

```
library(lattice)#For visualizations also required by caret
library(ggplot2)#For graphs also required by caret
```

```
#|eval: false
#Data initialization and preprocessing
data("Ionosphere")
summary(Ionosphere)
```

V1	V2	V3	V4	V5
0: 38	0:351	Min. : -1.0000	Min. : -1.00000	Min. : -1.0000
1:313		1st Qu.: 0.4721	1st Qu.: -0.06474	1st Qu.: 0.4127
		Median : 0.8711	Median : 0.01631	Median : 0.8092
		Mean : 0.6413	Mean : 0.04437	Mean : 0.6011
		3rd Qu.: 1.0000	3rd Qu.: 0.19418	3rd Qu.: 1.0000
		Max. : 1.0000	Max. : 1.00000	Max. : 1.0000
	V6	V7	V8	V9
	Min. : -1.0000	Min. : -1.0000	Min. : -1.00000	Min. : -1.00000
	1st Qu.: -0.0248	1st Qu.: 0.2113	1st Qu.: -0.05484	1st Qu.: 0.08711
	Median : 0.0228	Median : 0.7287	Median : 0.01471	Median : 0.68421
	Mean : 0.1159	Mean : 0.5501	Mean : 0.11936	Mean : 0.51185
	3rd Qu.: 0.3347	3rd Qu.: 0.9692	3rd Qu.: 0.44567	3rd Qu.: 0.95324
	Max. : 1.0000	Max. : 1.0000	Max. : 1.00000	Max. : 1.00000
	V10	V11	V12	V13
	Min. : -1.00000	Min. : -1.00000	Min. : -1.00000	Min. : -1.0000
	1st Qu.: -0.04807	1st Qu.: 0.02112	1st Qu.: -0.06527	1st Qu.: 0.0000
	Median : 0.01829	Median : 0.66798	Median : 0.02825	Median : 0.6441
	Mean : 0.18135	Mean : 0.47618	Mean : 0.15504	Mean : 0.4008
	3rd Qu.: 0.53419	3rd Qu.: 0.95790	3rd Qu.: 0.48237	3rd Qu.: 0.9555
	Max. : 1.00000	Max. : 1.00000	Max. : 1.00000	Max. : 1.0000
	V14	V15	V16	V17
	Min. : -1.00000	Min. : -1.0000	Min. : -1.00000	Min. : -1.0000
	1st Qu.: -0.07372	1st Qu.: 0.0000	1st Qu.: -0.08170	1st Qu.: 0.0000
	Median : 0.03027	Median : 0.6019	Median : 0.00000	Median : 0.5909
	Mean : 0.09341	Mean : 0.3442	Mean : 0.07113	Mean : 0.3819
	3rd Qu.: 0.37486	3rd Qu.: 0.9193	3rd Qu.: 0.30897	3rd Qu.: 0.9357
	Max. : 1.00000	Max. : 1.0000	Max. : 1.00000	Max. : 1.0000
	V18	V19	V20	V21
	Min. : -1.000000	Min. : -1.0000	Min. : -1.00000	Min. : -1.0000
	1st Qu.: -0.225690	1st Qu.: 0.0000	1st Qu.: -0.23467	1st Qu.: 0.0000
	Median : 0.000000	Median : 0.5762	Median : 0.00000	Median : 0.4991
	Mean : -0.003617	Mean : 0.3594	Mean : -0.02402	Mean : 0.3367
	3rd Qu.: 0.195285	3rd Qu.: 0.8993	3rd Qu.: 0.13437	3rd Qu.: 0.8949
	Max. : 1.000000	Max. : 1.0000	Max. : 1.00000	Max. : 1.0000
	V22	V23	V24	V25
	Min. : -1.000000	Min. : -1.0000	Min. : -1.00000	Min. : -1.0000

1st Qu.: -0.243870	1st Qu.: 0.0000	1st Qu.: -0.36689	1st Qu.: 0.0000
Median : 0.000000	Median : 0.5318	Median : 0.00000	Median : 0.5539
Mean : 0.008296	Mean : 0.3625	Mean : -0.05741	Mean : 0.3961
3rd Qu.: 0.188760	3rd Qu.: 0.9112	3rd Qu.: 0.16463	3rd Qu.: 0.9052
Max. : 1.000000	Max. : 1.0000	Max. : 1.00000	Max. : 1.0000
V26	V27	V28	V29
Min. : -1.00000	Min. : -1.0000	Min. : -1.00000	Min. : -1.0000
1st Qu.: -0.33239	1st Qu.: 0.2864	1st Qu.: -0.44316	1st Qu.: 0.0000
Median : -0.01505	Median : 0.7082	Median : -0.01769	Median : 0.4966
Mean : -0.07119	Mean : 0.5416	Mean : -0.06954	Mean : 0.3784
3rd Qu.: 0.15676	3rd Qu.: 0.9999	3rd Qu.: 0.15354	3rd Qu.: 0.8835
Max. : 1.00000	Max. : 1.0000	Max. : 1.00000	Max. : 1.0000
V30	V31	V32	V33
Min. : -1.00000	Min. : -1.0000	Min. : -1.000000	Min. : -1.0000
1st Qu.: -0.23689	1st Qu.: 0.0000	1st Qu.: -0.242595	1st Qu.: 0.0000
Median : 0.00000	Median : 0.4428	Median : 0.000000	Median : 0.4096
Mean : -0.02791	Mean : 0.3525	Mean : -0.003794	Mean : 0.3494
3rd Qu.: 0.15407	3rd Qu.: 0.8576	3rd Qu.: 0.200120	3rd Qu.: 0.8138
Max. : 1.00000	Max. : 1.0000	Max. : 1.000000	Max. : 1.0000
V34	Class		
Min. : -1.00000	bad : 126		
1st Qu.: -0.16535	good : 225		
Median : 0.00000			
Mean : 0.01448			
3rd Qu.: 0.17166			
Max. : 1.00000			

```
head(Ionosphere)
```

	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10
1	1	0	0.99539	-0.05889	0.85243	0.02306	0.83398	-0.37708	1.00000	0.03760
2	1	0	1.00000	-0.18829	0.93035	-0.36156	-0.10868	-0.93597	1.00000	-0.04549
3	1	0	1.00000	-0.03365	1.00000	0.00485	1.00000	-0.12062	0.88965	0.01198
4	1	0	1.00000	-0.45161	1.00000	1.00000	0.71216	-1.00000	0.00000	0.00000
5	1	0	1.00000	-0.02401	0.94140	0.06531	0.92106	-0.23255	0.77152	-0.16399
6	1	0	0.02337	-0.00592	-0.09924	-0.11949	-0.00763	-0.11824	0.14706	0.06637
	V11	V12	V13	V14	V15	V16	V17	V18		
1	0.85243	-0.17755	0.59755	-0.44945	0.60536	-0.38223	0.84356	-0.38542		
2	0.50874	-0.67743	0.34432	-0.69707	-0.51685	-0.97515	0.05499	-0.62237		
3	0.73082	0.05346	0.85443	0.00827	0.54591	0.00299	0.83775	-0.13644		
4	0.00000	0.00000	0.00000	0.00000	-1.00000	0.14516	0.54094	-0.39330		
5	0.52798	-0.20275	0.56409	-0.00712	0.34395	-0.27457	0.52940	-0.21780		

6	0.03786	-0.06302	0.00000	0.00000	-0.04572	-0.15540	-0.00343	-0.10196		
	V19	V20	V21	V22	V23	V24	V25	V26		
1	0.58212	-0.32192	0.56971	-0.29674	0.36946	-0.47357	0.56811	-0.51171		
2	0.33109	-1.00000	-0.13151	-0.45300	-0.18056	-0.35734	-0.20332	-0.26569		
3	0.75535	-0.08540	0.70887	-0.27502	0.43385	-0.12062	0.57528	-0.40220		
4	-1.00000	-0.54467	-0.69975	1.00000	0.00000	0.00000	1.00000	0.90695		
5	0.45107	-0.17813	0.05982	-0.35575	0.02309	-0.52879	0.03286	-0.65158		
6	-0.11575	-0.05414	0.01838	0.03669	0.01519	0.00888	0.03513	-0.01535		
	V27	V28	V29	V30	V31	V32	V33	V34	Class	
1	0.41078	-0.46168	0.21266	-0.34090	0.42267	-0.54487	0.18641	-0.45300	good	
2	-0.20468	-0.18401	-0.19040	-0.11593	-0.16626	-0.06288	-0.13738	-0.02447	bad	
3	0.58984	-0.22145	0.43100	-0.17365	0.60436	-0.24180	0.56045	-0.38238	good	
4	0.51613	1.00000	1.00000	-0.20099	0.25682	1.00000	-0.32382	1.00000	bad	
5	0.13290	-0.53206	0.02431	-0.62197	-0.05707	-0.59573	-0.04608	-0.65697	good	
6	-0.03240	0.09223	-0.07859	0.00732	0.00000	0.00000	-0.00039	0.12011	bad	

```
df <- Ionosphere
###About the data: 351 Observations and 34 Independent Variables(Removed one)
###Last column in the data is categorical variable called Class: good/bad
df <- subset(df, select = -V2) #Removed V2 bc all 0's
str(df)
```

```
'data.frame': 351 obs. of 34 variables:
 $ V1 : Factor w/ 2 levels "0","1": 2 2 2 2 2 2 2 1 2 2 ...
 $ V3 : num 0.995 1 1 1 1 ...
 $ V4 : num -0.0589 -0.1883 -0.0336 -0.4516 -0.024 ...
 $ V5 : num 0.852 0.93 1 1 0.941 ...
 $ V6 : num 0.02306 -0.36156 0.00485 1 0.06531 ...
 $ V7 : num 0.834 -0.109 1 0.712 0.921 ...
 $ V8 : num -0.377 -0.936 -0.121 -1 -0.233 ...
 $ V9 : num 1 1 0.89 0 0.772 ...
 $ V10 : num 0.0376 -0.0455 0.012 0 -0.164 ...
 $ V11 : num 0.852 0.509 0.731 0 0.528 ...
 $ V12 : num -0.1776 -0.6774 0.0535 0 -0.2028 ...
 $ V13 : num 0.598 0.344 0.854 0 0.564 ...
 $ V14 : num -0.44945 -0.69707 0.00827 0 -0.00712 ...
 $ V15 : num 0.605 -0.517 0.546 -1 0.344 ...
 $ V16 : num -0.38223 -0.97515 0.00299 0.14516 -0.27457 ...
 $ V17 : num 0.844 0.055 0.838 0.541 0.529 ...
 $ V18 : num -0.385 -0.622 -0.136 -0.393 -0.218 ...
 $ V19 : num 0.582 0.331 0.755 -1 0.451 ...
 $ V20 : num -0.3219 -1 -0.0854 -0.5447 -0.1781 ...
```

```


$ V21 : num  0.5697 -0.1315 0.7089 -0.6997 0.0598 ...
$ V22 : num  -0.297 -0.453 -0.275 1 -0.356 ...
$ V23 : num  0.3695 -0.1806 0.4339 0 0.0231 ...
$ V24 : num  -0.474 -0.357 -0.121 0 -0.529 ...
$ V25 : num  0.5681 -0.2033 0.5753 1 0.0329 ...
$ V26 : num  -0.512 -0.266 -0.402 0.907 -0.652 ...
$ V27 : num  0.411 -0.205 0.59 0.516 0.133 ...
$ V28 : num  -0.462 -0.184 -0.221 1 -0.532 ...
$ V29 : num  0.2127 -0.1904 0.431 1 0.0243 ...
$ V30 : num  -0.341 -0.116 -0.174 -0.201 -0.622 ...
$ V31 : num  0.4227 -0.1663 0.6044 0.2568 -0.0571 ...
$ V32 : num  -0.5449 -0.0629 -0.2418 1 -0.5957 ...
$ V33 : num  0.1864 -0.1374 0.5605 -0.3238 -0.0461 ...
$ V34 : num  -0.453 -0.0245 -0.3824 1 -0.657 ...
$ Class: Factor w/ 2 levels "bad","good": 2 1 2 1 2 1 2 1 2 1 ...

```

```

#Here are some graphical summaries of the Ionosphere data
hist(df$V7, main = "Histogram of V7 Occurance")


```




```

hist(df$V10, main = "Histogram of V10 Occurance")

```




```
boxplot(V7 ~ Class, data = df, col = c("blue", "green"), main = "Boxplot of V7 by Class")
```





```
boxplot(V10 ~ Class, data = df, col = c("blue", "green"), main = "Boxplot of V10 by Class")
```



```
#Here are the corresponding numerical summaries
summary(df$V7[df$Class == "good"])
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-0.2222	0.5178	0.8159	0.7159	0.9603	1.0000

```
summary(df$V7[df$Class == "bad"])
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-1.000000	0.000000	0.007185	0.253984	1.000000	1.000000

```
summary(df$V10[df$Class == "good"])
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
-0.51171	-0.04286	0.04317	0.22496	0.56830	1.00000

```
summary(df$V10[df$Class == "bad"])
```

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
	-1.00000	-0.06653	0.00000	0.10346	0.50075	1.00000

There is some skewness in the histogram of V7 to the left direction.

As for the box plot of V10 the two plots seem to be roughly identical and this follows its histogram which looks approximately normally distributed.

```
set.seed(4323)
#Data slicing
intrainQ3 <- createDataPartition(y = df$Class, p= 0.7, list = FALSE)
trainingQ3 <- df[intrainQ3,]
testingQ3 <- df[-intrainQ3,]

#checking to see if our dimensions add up
dim(trainingQ3)
```

```
[1] 247  34
```

```
dim(testingQ3)
```

```
[1] 104  34
```

```
#
trControl <- trainControl(method = "cv",
                          number = 5)
fit <- train(Class ~ .,
            method = "knn",
            trControl = trControl,
            tuneGrid = expand.grid(k = 1:10),
            data = df)
fit
```

k-Nearest Neighbors

351 samples  
33 predictor

2 classes: 'bad', 'good'

No pre-processing

Resampling: Cross-Validated (5 fold)

Summary of sample sizes: 281, 280, 281, 281, 281

Resampling results across tuning parameters:

k	Accuracy	Kappa
1	0.8461569	0.6409977
2	0.8375855	0.6208973
3	0.8375855	0.6155820
4	0.8461167	0.6360559
5	0.8462374	0.6359032
6	0.8348491	0.6049333
7	0.8319920	0.6004563
8	0.8206036	0.5691947
9	0.8320322	0.6002482
10	0.8263179	0.5847849

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was k = 5.

```
test_predictionQ3 <- predict(fit, newdata= testingQ3)
test_predictionQ3
```

```
[1] good good good bad good bad good good good good good good good good good
[16] good bad good bad good bad bad good good good good good good bad good good
[31] good bad bad bad good good good good bad good good good bad good good
[46] bad good good bad good good bad good bad bad good bad good good bad
[61] good good good good bad good good bad good bad good bad good good good
[76] good good good good good good good good bad good good good good good good
[91] good good good good good good good good good good good good good good
Levels: bad good
```

Test error for k = 5 was found to be the best at roughly 15.38% in comparison to k = 7, 16.8% and k = 1, 15.39%.

```
confusionMatrix(test_predictionQ3, testingQ3$Class)
```

Confusion Matrix and Statistics

```

      Reference
Prediction bad good
      bad   22    2
      good   15   65

      Accuracy : 0.8365
      95% CI : (0.7512, 0.9018)
No Information Rate : 0.6442
P-Value [Acc > NIR] : 1.185e-05

```

Kappa : 0.613

McNemar's Test P-Value : 0.003609

```

      Sensitivity : 0.5946
      Specificity : 0.9701
Pos Pred Value : 0.9167
Neg Pred Value : 0.8125
      Prevalence : 0.3558
      Detection Rate : 0.2115
Detection Prevalence : 0.2308
      Balanced Accuracy : 0.7824

```

'Positive' Class : bad

What is the confusion matrix saying?

There is an 83.65% accuracy in this model.

The model hallucinated 15 bad as good and 2 good as bad, in which this model is better at determining what is good opposed to bad.

#### Question 4:

```

#Data initialization and preprocessing
data("Auto")
summary(Auto)

```

```

      mpg      cylinders      displacement      horsepower      weight
Min.   : 9.00   Min.   :3.000   Min.   : 68.0   Min.   : 46.0   Min.   :1613

```

1st Qu.:17.00	1st Qu.:4.000	1st Qu.:105.0	1st Qu.: 75.0	1st Qu.:2225
Median :22.75	Median :4.000	Median :151.0	Median : 93.5	Median :2804
Mean :23.45	Mean :5.472	Mean :194.4	Mean :104.5	Mean :2978
3rd Qu.:29.00	3rd Qu.:8.000	3rd Qu.:275.8	3rd Qu.:126.0	3rd Qu.:3615
Max. :46.60	Max. :8.000	Max. :455.0	Max. :230.0	Max. :5140

acceleration	year	origin		name
Min. : 8.00	Min. :70.00	Min. :1.000	amc matador	: 5
1st Qu.:13.78	1st Qu.:73.00	1st Qu.:1.000	ford pinto	: 5
Median :15.50	Median :76.00	Median :1.000	toyota corolla	: 5
Mean :15.54	Mean :75.98	Mean :1.577	amc gremlin	: 4
3rd Qu.:17.02	3rd Qu.:79.00	3rd Qu.:2.000	amc hornet	: 4
Max. :24.80	Max. :82.00	Max. :3.000	chevrolet chevette:	4
			(Other)	:365

```
attach(Auto)
```


The following object is masked from package:ggplot2:

```
mpg
```

```
mpg01 <- ifelse( mpg > median(mpg), yes = 1, no = 0)
newAuto <- data.frame(Auto, mpg01)
str(newAuto)
```

```
'data.frame': 392 obs. of 10 variables:
 $ mpg      : num  18 15 18 16 17 15 14 14 14 15 ...
 $ cylinders : num  8 8 8 8 8 8 8 8 8 8 ...
 $ displacement: num  307 350 318 304 302 429 454 440 455 390 ...
 $ horsepower  : num  130 165 150 150 140 198 220 215 225 190 ...
 $ weight      : num  3504 3693 3436 3433 3449 ...
 $ acceleration: num  12 11.5 11 12 10.5 10 9 8.5 10 8.5 ...
 $ year        : num  70 70 70 70 70 70 70 70 70 70 ...
 $ origin      : num  1 1 1 1 1 1 1 1 1 1 ...
 $ name        : Factor w/ 304 levels "amc ambassador brougham",...: 49 36 231 14 161 141 54 ...
 $ mpg01       : num  0 0 0 0 0 0 0 0 0 0 ...
```

```
#Scatter plot matrix
pairs(newAuto)
```



Scatter plot matrix:

Of the variables in the data set Auto, I found cylinders, weight, displacement, horsepower, acceleration and the age of the car to be among the most influencing of Mpg01.

```
#Here we standardized the data, since knn works on distance we don't want anything skewed
#cbind makes a matrix, apply, applies the scale (standardize) to the columns 2 not the rows
newAuto <- data.frame(mpg01, apply(cbind(cylinders, weight, displacement, horsepower, acceleration, year), 2, FUN=scale))
str(newAuto)
```

```
'data.frame':  392 obs. of  7 variables:
 $ mpg01      : num  0 0 0 0 0 0 0 0 0 0 ...
 $ cylinders   : num  1.48 1.48 1.48 1.48 1.48 ...
 $ weight      : num  0.62 0.842 0.54 0.536 0.555 ...
 $ displacement: num  1.08 1.49 1.18 1.05 1.03 ...
 $ horsepower  : num  0.663 1.573 1.183 1.183 0.923 ...
 $ acceleration: num -1.28 -1.46 -1.65 -1.28 -1.83 ...
 $ year        : num  70 70 70 70 70 70 70 70 70 70 ...
```

```
#Splitting the data 70:30
intrainQ4 <- createDataPartition(y = mpg01, p= 0.7, list = FALSE)
trainingQ4 <- newAuto[intrainQ4,]
testingQ4 <- newAuto[-intrainQ4,]
```

```
dim(trainingQ4)
```

```
[1] 276  7
```

```
dim(testingQ4)
```

```
[1] 116  7
```

```
set.seed(1)

trControlQ4 <- trainControl(method = "cv",
                             number = 5)
fitQ4 <- train(as.factor(mpg01) ~ .,
               method      = "knn",
               trControl    = trControlQ4,
               tuneGrid     = expand.grid(k = 1:10),
               data         = newAuto)
fitQ4
```

k-Nearest Neighbors

392 samples  
6 predictor  
2 classes: '0', '1'

No pre-processing  
Resampling: Cross-Validated (5 fold)  
Summary of sample sizes: 314, 313, 314, 313, 314  
Resampling results across tuning parameters:

k	Accuracy	Kappa
1	0.9209672	0.8419133
2	0.9210321	0.8420609
3	0.9260954	0.8522152
4	0.9108731	0.8217583
5	0.9133723	0.8268230
6	0.9159364	0.8319125
7	0.9185329	0.8371063
8	0.9185005	0.8370503

```

9 0.9236611 0.8473627
10 0.9236611 0.8473529

```

Accuracy was used to select the optimal model using the largest value.  
The final value used for the model was  $k = 3$ .

```

test_predictionQ4 <- predict(fitQ4, newdata= testingQ4)
test_predictionQ4

```

```

[1] 0 0 0 1 0 0 1 1 1 0 0 0 0 0 1 1 1 1 0 0 0 0 0 0 1 0 0 0 0 0 0 1 0 0 1 0
[38] 1 0 0 1 1 0 0 0 0 0 0 1 1 1 1 1 0 0 0 0 0 1 1 1 1 0 1 0 1 0 0 1 1 0 1 0 0
[75] 1 0 0 0 1 1 1 1 0 0 0 0 0 1 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
[112] 1 1 1 1 1
Levels: 0 1

```

Among the best K-Values,  $k = 3$  was found to have a  $(1-0.926)$  7.4% error rate.

I unfortunately already scaled my data but here is why: Because KNN works on distance, certain variables will have disproportionate weight. We needed to standardize the variables so that the distance between data points is not skewed. So scaling each feature to have a mean of 0 and a standard deviation of 1 should help.

```

confusionMatrix(test_predictionQ4, as.factor(testingQ4$mpg01))

```

#### Confusion Matrix and Statistics

```

          Reference
Prediction 0  1
0  54  1
1  4  57

      Accuracy : 0.9569
      95% CI   : (0.9023, 0.9859)
No Information Rate : 0.5
P-Value [Acc > NIR] : <2e-16

      Kappa   : 0.9138

McNemar's Test P-Value : 0.3711

      Sensitivity : 0.9310

```



```

        Specificity : 0.9828
        Pos Pred Value : 0.9818
        Neg Pred Value : 0.9344
        Prevalence : 0.5000
        Detection Rate : 0.4655
        Detection Prevalence : 0.4741
        Balanced Accuracy : 0.9569

```

```
'Positive' Class : 0
```

### Question 5:

```
head(Auto)
```

```

      mpg cylinders displacement horsepower weight acceleration year origin
1   18         8         307         130   3504          12.0    70      1
2   15         8         350         165   3693          11.5    70      1
3   18         8         318         150   3436          11.0    70      1
4   16         8         304         150   3433          12.0    70      1
5   17         8         302         140   3449          10.5    70      1
6   15         8         429         198   4341          10.0    70      1
      name
1 chevrolet chevelle malibu
2      buick skylark 320
3    plymouth satellite
4      amc rebel sst
5      ford torino
6    ford galaxie 500

```

```
summary(Auto)
```

```

      mpg      cylinders displacement horsepower      weight
Min.   : 9.00   Min.   :3.000   Min.   : 68.0   Min.   : 46.0   Min.   :1613
1st Qu.:17.00   1st Qu.:4.000   1st Qu.:105.0   1st Qu.: 75.0   1st Qu.:2225
Median :22.75   Median :4.000   Median :151.0   Median : 93.5   Median :2804
Mean   :23.45   Mean   :5.472   Mean   :194.4   Mean   :104.5   Mean   :2978
3rd Qu.:29.00   3rd Qu.:8.000   3rd Qu.:275.8   3rd Qu.:126.0   3rd Qu.:3615
Max.   :46.60   Max.   :8.000   Max.   :455.0   Max.   :230.0   Max.   :5140

```

acceleration	year	origin	name	
Min. : 8.00	Min. :70.00	Min. :1.000	amc matador	: 5
1st Qu.:13.78	1st Qu.:73.00	1st Qu.:1.000	ford pinto	: 5
Median :15.50	Median :76.00	Median :1.000	toyota corolla	: 5
Mean :15.54	Mean :75.98	Mean :1.577	amc gremlin	: 4
3rd Qu.:17.02	3rd Qu.:79.00	3rd Qu.:2.000	amc hornet	: 4
Max. :24.80	Max. :82.00	Max. :3.000	chevrolet chevette:	4
			(Other)	:365

```
set.seed(1)
```

```
trControlQ5 <- trainControl(method = "LOOCV", number = 5)
newAutoQ5 <- data.frame(Auto, mpg01)

fitQ5 <- train(as.factor(mpg01) ~ .,
               method = "knn",
               trControl = trControlQ5,
               tuneGrid = expand.grid(k = 1:10),
               data = newAutoQ5)

fitQ5
```

k-Nearest Neighbors

392 samples

9 predictor

2 classes: '0', '1'

No pre-processing

Resampling: Leave-One-Out Cross-Validation

Summary of sample sizes: 391, 391, 391, 391, 391, 391, ...

Resampling results across tuning parameters:

k	Accuracy	Kappa
1	0.8698980	0.7397959
2	0.8724490	0.7448980
3	0.8877551	0.7755102
4	0.8801020	0.7602041
5	0.8826531	0.7653061
6	0.8750000	0.7500000
7	0.8724490	0.7448980
8	0.8877551	0.7755102
9	0.8750000	0.7500000

```
10 0.8673469 0.7346939
```

Accuracy was used to select the optimal model using the largest value.  
The final value used for the model was  $k = 8$ .

Before the scaling, the approach towards validation that has proven to be better is a scaled k-cross validation approach. The accuracy of the best performing k values from each approach were ascertain, and k-cross validation prevailed, with a test error rate of  $(1 - .926) = 7.4\%$  in comparison to a non scaled LOOCV test error rate where  $k = 8$  of  $(1 - .888) 11.2\%$

```
set.seed(1)
newAutoQ5Scaled <- data.frame(mpg01, apply(cbind(cylinders, weight, displacement, horsepower),
fitQ5 <- train(as.factor(mpg01) ~ .,
               method      = "knn",
               trControl    = trControlQ5,
               tuneGrid     = expand.grid(k = 1:10),
               data         = newAutoQ5Scaled)
fitQ5
```

k-Nearest Neighbors

```
392 samples
 6 predictor
 2 classes: '0', '1'
```

No pre-processing

Resampling: Leave-One-Out Cross-Validation

Summary of sample sizes: 391, 391, 391, 391, 391, 391, ...

Resampling results across tuning parameters:

k	Accuracy	Kappa
1	0.9234694	0.8469388
2	0.9081633	0.8163265
3	0.9209184	0.8418367
4	0.9107143	0.8214286
5	0.9132653	0.8265306
6	0.9158163	0.8316327
7	0.9158163	0.8316327
8	0.9234694	0.8469388
9	0.9234694	0.8469388
10	0.9183673	0.8367347

Accuracy was used to select the optimal model using the largest value.  
The final value used for the model was  $k = 9$ .

After having scaled the data it seems that LOOCV performed better than its un-scaled version with a  $k$  value = 9 yielding a test error rate of (1-.9235) 7.7% but exceptionally close to a scaled version of  $k$  - cross validation with a difference of approximately 0.3%.

In this case it would be better to do the  $k$  - cross validation approach as there is less compute time, as far as trusting which validation approach more, that would also be  $k$ -cross validation but it depends on the use case of your machine learning model.

PS: I couldnt figure out how to display the code but prevent evaluation, I am sorry.